Network Architecture: The model consists of 3 hidden layers, each using the ReLU activation function, and a final sigmoid activation layer to handle binary classification.

Loss Function: Binary Cross-Entropy Loss was used for optimizing the model's predictions.

Optimizer: The Adam optimizer was selected, with a learning rate set to 0.001 to facilitate efficient training.

Evaluation Metrics: The model's performance was evaluated using Accuracy, Precision, Recall, and F1-Score metrics.

Challenges: Managing potential overfitting or underfitting required adjusting hyperparameters and considering regularization techniques.